

GOPAL PUJARI

Azure Data Engineer

Bangalore, India | +91 90328 91455 | gopalpujari.de@gmail.com | [LinkedIn](#)

PROFESSIONAL SUMMARY

Azure Data Engineer with 4.6 years of experience designing and delivering ETL/ELT pipelines, data models, and Delta Lake-based data warehousing solutions across e-commerce, healthcare, construction finance, and energy. Builds production-grade pipelines using Azure Data Factory, Azure Databricks (Unity Catalog), PySpark, and Spark SQL, meaningfully reducing ETL runtime through incremental loading and Spark performance tuning. Experienced in Medallion (bronze/silver/gold) data modeling, HIPAA-compliant data handling, and Azure SQL-based warehousing, with a strong background in SQL Server, MySQL performance tuning, CI/CD delivery, and Agile practices. Proven track record mentoring junior engineers and partnering with cross-functional and client teams to deliver trusted, production-ready datasets.

CORE SKILLS

- Programming: Python, SQL
- Big Data & Processing: Apache Spark, PySpark, Spark SQL, Delta Lake
- Data Modeling & Warehousing: Dimensional modeling, Medallion architecture (Bronze/Silver/Gold), Azure SQL Data Warehouse
- Orchestration / ETL: Azure Data Factory (pipelines, linked services, datasets, triggers)
- Azure Cloud: Azure Databricks, Unity Catalog, ADLS Gen2, Azure Blob Storage, Azure SQL Database, Azure Key Vault, Azure Monitor
- Databases: MS SQL Server, MySQL
- File Formats: Parquet, Avro, JSON, CSV, Delta
- DevOps & Tools: Git, GitHub, Azure DevOps, CI/CD, JIRA, PyCharm, VS Code, Jupyter
- Reporting: Power BI

PROFESSIONAL EXPERIENCE

Flipkart

Azure Data Engineer

Project: E-Commerce | Client: E-Kart Logistics

Feb 2026 – Present

Bangalore, India

- Architect metadata-driven ETL pipelines in Azure Data Factory, ingesting data from multiple enterprise databases, REST APIs, and flat files into ADLS Gen2, standardizing and scaling multi-source data ingestion.
- Design PySpark notebooks in Azure Databricks (Unity Catalog-governed) to cleanse, transform, and process high-volume transactional and customer datasets using Spark DataFrame APIs and Spark SQL, handling large record volumes daily.
- Implement Bronze-Silver-Gold Medallion architecture on Delta Lake, incorporating watermark-based incremental loading and MERGE operations to optimize data processing and maintain reliable, auditable pipelines.
- Optimize Apache Spark workloads through effective partitioning, caching, and broadcast joins, meaningfully improving pipeline throughput and reducing overall ETL runtime.
- Implement robust error handling, pipeline logging, monitoring, and source-to-target data validation using Azure Monitor, maintaining strong SLA compliance and consistent data delivery.
- Own CI/CD pipeline delivery for data workflows via Azure DevOps and Git, including secrets and credential management through Azure Key Vault, following Agile development practices.
- Deliver curated and trusted datasets for Power BI reports by collaborating closely with cross-functional analytics and business teams.
- Mentor junior engineers through code reviews and pair programming on PySpark and Azure Data Factory best practices, authoring technical documentation for pipeline standards reused across teams.

DEC Industries Pvt Ltd

Data Engineer

Project: Construction Finance Management System | Client: NDRS

Apr 2025 – Feb 2026

Hyderabad, India

- Engineered and maintained ETL pipelines to orchestrate data workflows across finance, construction, and operational source systems.
- Automated job scheduling, workflow orchestration, and dependency management for batch and near-real-time data processing, minimizing manual intervention and improving operational efficiency.
- Implemented robust error handling, retry mechanisms, and alerting to improve pipeline reliability, fault tolerance, and overall system availability, cutting pipeline downtime.
- Optimized workflow performance by parallelizing independent tasks and refining pipeline execution strategies, reducing processing overhead and execution time.
- Established end-to-end monitoring using Azure Monitor and developed custom dashboards to track pipeline health, execution status, failures, and system performance.
- Delivered trusted financial datasets to support cost tracking and forecasting while collaborating with cross-functional teams across finance, construction, and operations.

Tech Mahindra

Data Engineer

Jul 2024 – Apr 2025

Bengaluru, India

Project: Healthcare Data Modernization System | Client: HCI Group

- Partnered with HCI Group on healthcare IT modernization initiatives spanning payer, provider, and patient data domains.
- Developed ETL pipelines using Azure Data Factory and Azure Databricks to integrate and standardize claims, hospital, and patient data from diverse source systems.
- Developed PySpark and SQL-based data processing solutions with a focus on HIPAA-compliant data handling, secure access via Azure Key Vault, interoperability, and compliance requirements.
- Automated claims processing and patient analytics workflows, meaningfully reducing manual processing effort and improving operational efficiency.
- Designed and implemented scalable data models to support analytical workloads and predictive insights into patient outcomes and healthcare operations.
- Migrated legacy healthcare applications and data workloads to Microsoft Azure, improving platform reliability, scalability, and infrastructure efficiency, and lowering costs.

Wipro Ltd.

Project Engineer

Jan 2022 – Jul 2024

Bangalore, India

Project: Gridstar – Smart Energy & Grid Financial Solutions | Client: Hitachi Energy

- Engineered end-to-end data solutions for Gridstar, a smart energy management platform supporting grid performance monitoring and financial planning.
- Developed ETL pipelines using Python, SQL, and Azure Data Factory to ingest grid and cost data from SFTP and SQL Server into Azure Data Lake Storage Gen2 (ADLS Gen2).
- Developed PySpark transformations in Azure Databricks to integrate construction finance and grid data, enabling comprehensive cost tracking and forecasting.
- Converted and standardized multi-format raw data, including CSV and JSON, into optimized Delta Lake tables, improving data accessibility and analytical query performance.
- Implemented incremental data processing using PySpark and SQL, enabling timely updates to grid and financial datasets while optimizing pipeline efficiency.
- Delivered integrated grid and financial analytics datasets to support data-driven resource allocation and operational decision-making across stakeholder teams.

PROJECTS

- Medallion Architecture Pipeline — Databricks / Unity Catalog
 - Built an end-to-end Bronze-Silver-Gold pipeline on a relational sales dataset (classicModels schema) under Unity Catalog governance, resolving schema mismatches, type errors, and CSV parsing issues during ingestion.
 - Applied DROP TABLE IF EXISTS patterns to manage schema evolution across Delta tables, validating each medallion layer for data quality and consistency through to analytics-ready gold tables.

EDUCATION

Bachelor of Technology

Jul 2021

Jawaharlal Nehru Technological University, Andhra Pradesh